

BENCHMARKING

DESIGN DOCUMENT

Shalin A. Chopra

2/12/2016

DESIGN DOCUMENT

1. CPU BENCHMARKING:

- The code for CPU benchmarking is written in C programming language.
- This program calculates the Integer and Floating point operations, in terms of GIOPS and GFLOPS. The main aim is to utilize the complete CPU cycles by executing different arithmetic instructions. Utilizing the CPU's Floating point unit (FPU) completely so that it gives us the maximum FLOPS.
- The program has performs 17 operations which run for 1 Billion times in a loop. The instructions like addition, multiplication, subtraction and division are performed, condition check are accounted as these operations. The function intBM() calculates the GIOPS and floatBM() calculates the GFLOPS.
- There are a maximum of 4 threads, that are executed for each of these functions. All the results are automated using scripts, the final output is directly displayed onto the screen.
- There is a 600 sec performance evaluation, which evaluates the IOPS and FLOPS executed per second, using 4 threads. The entire results are automated the user directly see two files which has values for each second.

2. DISK BENCHMARKING:

- The Disk benchmarking operation is implemented using JAVA programming language.
- The design includes implementation for 3 different block sizes i.e. 1B, 1KB and 1MB each for Sequential and Random operations.
- It implements 4 methods, sequential read & write and random read & write.
- The sequential access is done using a file, and data is read from the file and written into it, in a sequential manner.
- For random access, a random number is generated which lies within the file size, and is seeked to that location onto that file and read and write operations are performed.
- There are 1, 2 threads used to account for concurrency, which helps in calculating Throughput and Latency.

3. MEMORY BENCHMARKING:

- This code is written in C programming language.
- For different block sizes i.e. 1 Byte, 1 KB and 1 MB, sequential and random access to the memory is made.

- This code also implements multi threading to achieve concurrency.
- The disk access are made using memcpy() function which is used to perform read and write operation onto the memory.
- This experiment was run on AWS thus not enough memory was present, the program allocated 50 MB of memory using malloc() function. Larger memory resulted in Segmentation Faults.
- The program also calculates the throughput and latency and helps in finding the Speeds of memory access.

4. NETWORK BENCHMARKING:

- This benchmark has been implemented using JAVA programming language.
- The benchmarking is done for both TCP as well as UDP protocol. The code is written to be executed on two different instances of AWS.
- This code does the basic packets transmission from Client to Server and back again, while implementing this we find the RTT of the transmission.
- The packets transmitted are of various sizes i.e. 1 Byte, 1KB and 1MB.
- The TCP being reliable and connection oriented requires pre connection setup and accepting of connection between client and server.
- On the other hand UDP being connection less, the packets are sent and received without and pre established connection.
- The benchmarking helps to find the Network Bandwidth and Latency.

Design Trade-offs made:

CPU Benchmark:

Can be made more efficient and utilize the complete cycles, by implementing more complex instruction such as LINPACK which uses Matrix Multiplication and Linear equation solving. So, that more FLOPS/IOPS can be achieved.

Disk Benchmark:

Since this code is implemented and executed on t2.micro instance which is a single core instance and doesn't always provide with all the resources available to them as these are shared resources, the data transfer rate is limited by the core. For, future implementation we can extend the design to multi- core environment and achieve better efficiency.

Memory Benchmark:

The memory size was limited and hence most part of it wasn't accessible to run the benchmark. The benchmark is implemented for 50 MB memory allocation. More memory can be allocated and the cache's can be isolated so as to achieve better performance.

BENCHMARKING

DESIGN DOCUMENT

Shalin A. Chopra

2/12/2016

PERFORMANCE EVALUATION

This document presents with the performance evaluation for 4 different Benchmarks.

I have evaluated the following benchmarks:


1. CPU
2. DISK
3. MEMORY
4. NETWORK

1. CPU BENCHMARKING

Performance is done on AWS EC2 instance i.e. t2.micro, linux version: Ubuntu
The specifications are as follows:

SPECIFICATIONS:	
Model Name	Intel(R) Xeon(R) CPU E5-2670 v2 @ 2.50GHz
CPU MHz	2500.056
Cache Size	25600 KB
CPU Cores (vCore)	1
RAM	1 GB
Memory	8 GB

The graphs & tables below show the performance of benchmarking:



The above graph depicts the Average GIOPS and GFLOPS for 1, 2, 4 threads.
GIOPS for Integer operation and GFLOPS for Floating point operations.

As seen in the graph, as the number of thread increased the Number of Instruction per cycle started to decrease. Maximum FLOPS and IOPS were obtained for 1 thread. If we run the experiment for more threads, we see the trend that the Instructions per cycle start to stabilize at a point. It is not always the case that, more the number of threads the better the performance is.

Average & Standard Deviation for 3 Experiments:

Average:


No. of Threads	Average GIOPS	Average GFLOPS
1	6.597515	6.184995
2	6.449359	6.029395
4	6.441324	6.060415

Standard Deviation:


No. of Threads	Standard Deviation
1	0.130069
2	0.048843
4	0.01059

Plotted Values for 600 sec experiment:

Integer Operations per second:



Floating point Operations per second:



The above graphs depicts the number of instructions (operations) executed per second, for 600 seconds having 4 threads running concurrently. We observe that there is a variation for every second, the number of operations doesn't move at a constant rate. This might be due to the CPU prioritization. The **Trendline** in (red) shows what expected result was.

Theoretical Performance:

Number of Cores * Number of Instruction per cycle * Clock Speed = $1 * 4 * 2.5$
= 10 GFLOPS

Compared to theoretical performance efficiency achieved is around = 62.5%

LINPACK BENCHMARK:

I ran the LINPACK Benchmark, at two different time intervals, having different values. For 1st the maximal value obtained is: **36.9150 GFLOPS**

For 2nd the maximal value obtained is: **20.778 GFLOPS**

This variation might be due to AWS's CPU priority and utilization as amazon doesn't provide with 100% CPU utilization of the instances as these are shared, it provides only around 10% utilization.

Intel(R) Optimized LINPACK Benchmark data

Current date/time: Sat Feb 13 01:33:58 2016

CPU frequency: 2.825 GHz

Number of CPUs: 1

Number of cores: 1

Number of threads: 1

Parameters are set to:

Number of tests: 15

Number of equations to solve (problem size)	: 1000	2000	5000	10000	15000	18000	20000	22000	25000	26000	27000	30000	35000	40000	45000
Leading dimension of array	: 1000	2000	5008	10000	15000	18008	20016	22008	25000	26000	27000	30000	35000	40000	45000
Number of trials to run	: 4	2	2	2	2	2	2	2	2	2	1	1	1	1	1
Data alignment value (in Kbytes)	: 4	4	4	4	4	4	4	4	4	4	4	1	1	1	1

Maximum memory requested that can be used=800204096, at the size=10000

===== Timing linear equation system solver =====

Size	LDA	Align.	Time(s)	GFlops	Residual	Residual(norm)	Check
1000	1000	4	0.037	17.8627	9.632295e-13	3.284860e-02	pass
1000	1000	4	0.025	27.0416	9.632295e-13	3.284860e-02	pass
1000	1000	4	0.025	27.1198	9.632295e-13	3.284860e-02	pass
1000	1000	4	0.025	26.8425	9.632295e-13	3.284860e-02	pass
2000	2000	4	0.187	28.6389	4.746648e-12	4.129002e-02	pass
2000	2000	4	0.187	28.6115	4.746648e-12	4.129002e-02	pass
5000	5008	4	2.455	33.9686	2.651185e-11	3.696863e-02	pass
5000	5008	4	2.439	34.1811	2.651185e-11	3.696863e-02	pass
10000	10000	4	18.143	36.7564	9.014595e-11	3.178637e-02	pass
10000	10000	4	18.065	36.9150	9.014595e-11	3.178637e-02	pass

Performance Summary (GFlops)

Size	LDA	Align.	Average	Maximal
1000	1000	4	24.7167	27.1198
2000	2000	4	28.6252	28.6389
5000	5008	4	34.0749	34.1811
10000	10000	4	36.8357	36.9150

Residual checks PASSED

End of tests

the correct number of CPUs/threads, problem input files, etc..

./runme_xeon64: 37: [: -gt: unexpected operator

Sat Feb 13 03:35:29 UTC 2016

Intel(R) Optimized LINPACK Benchmark data

Current date/time: Sat Feb 13 03:35:29 2016

CPU frequency: 2.929 GHz

Number of CPUs: 1

Number of cores: 1

Number of threads: 1

Parameters are set to:

Number of tests: 15

Number of equations to solve (problem size)	: 1000	2000	5000	10000	15000	18000	20000	22000	25000	26000	27000	30000	35000	40000	45000
Leading dimension of array	: 1000	2000	5008	10000	15000	18008	20016	22008	25000	26000	27000	30000	35000	40000	45000
Number of trials to run	: 4	2	2	2	2	2	2	2	2	2	1	1	1	1	1
Data alignment value (in Kbytes)	: 4	4	4	4	4	4	4	4	4	4	4	1	1	1	1

Maximum memory requested that can be used=800204096, at the size=10000

===== Timing linear equation system solver =====

Size	LDA	Align.	Time(s)	GFlops	Residual	Residual(norm)	Check
1000	1000	4	0.051	13.1048	9.900691e-13	3.376390e-02	pass
1000	1000	4	0.038	17.4613	9.900691e-13	3.376390e-02	pass
1000	1000	4	0.038	17.6799	9.900691e-13	3.376390e-02	pass
1000	1000	4	0.038	17.7428	9.900691e-13	3.376390e-02	pass
2000	2000	4	0.286	18.6734	4.053480e-12	3.526031e-02	pass
2000	2000	4	0.280	19.0483	4.053480e-12	3.526031e-02	pass
5000	5008	4	4.140	20.1390	2.336047e-11	3.257429e-02	pass
5000	5008	4	4.113	20.2752	2.336047e-11	3.257429e-02	pass
10000	10000	4	32.187	20.7185	1.124127e-10	3.963786e-02	pass
10000	10000	4	32.106	20.7708	1.124127e-10	3.963786e-02	pass

Performance Summary (GFlops)

Size	LDA	Align.	Average	Maximal
1000	1000	4	16.4972	17.7428
2000	2000	4	18.8609	19.0483
5000	5008	4	20.2071	20.2752
10000	10000	4	20.7446	20.7708

Residual checks PASSED

End of tests

Done: Sat Feb 13 03:36:49 UTC 2016


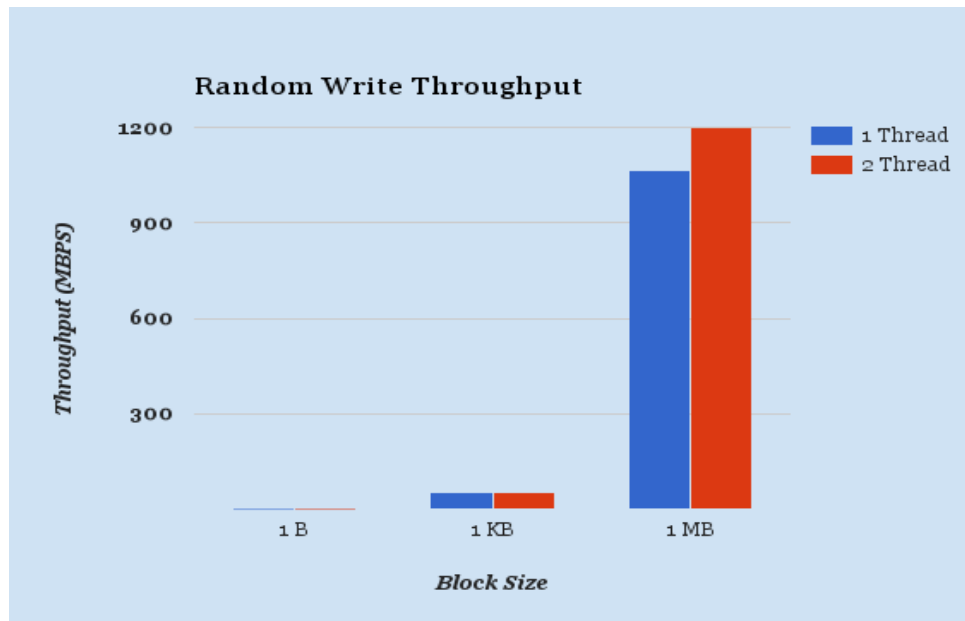
ubuntu@ip-172-31-25-167:~/l_mklb_p_11.3.1.002/benchmarks_11.3.1/linux/mkl/benchmarks/linpack\$

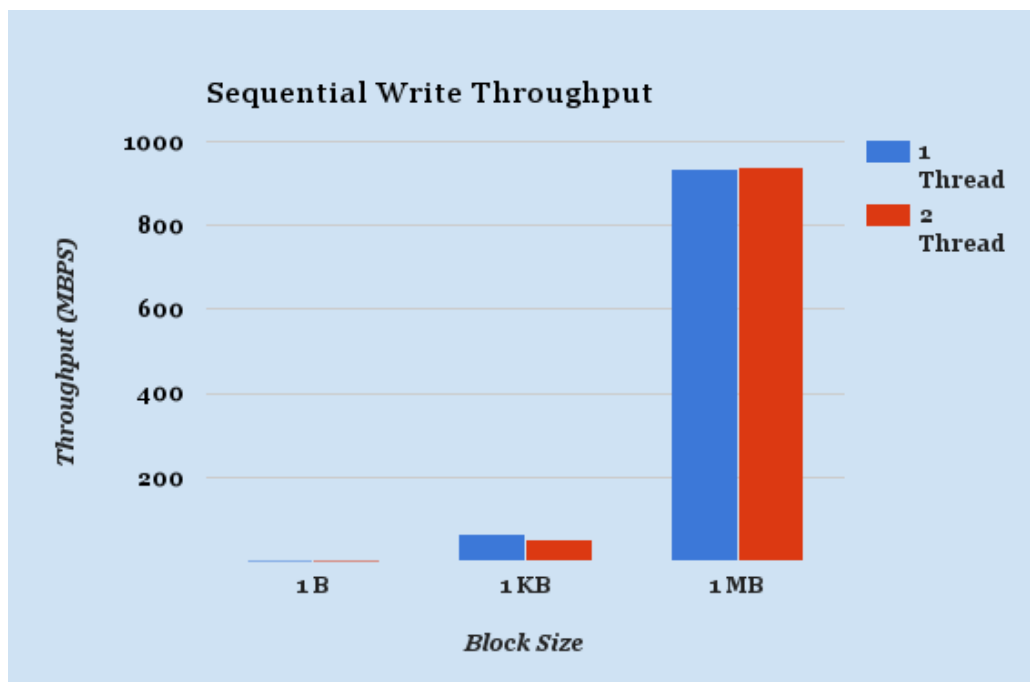
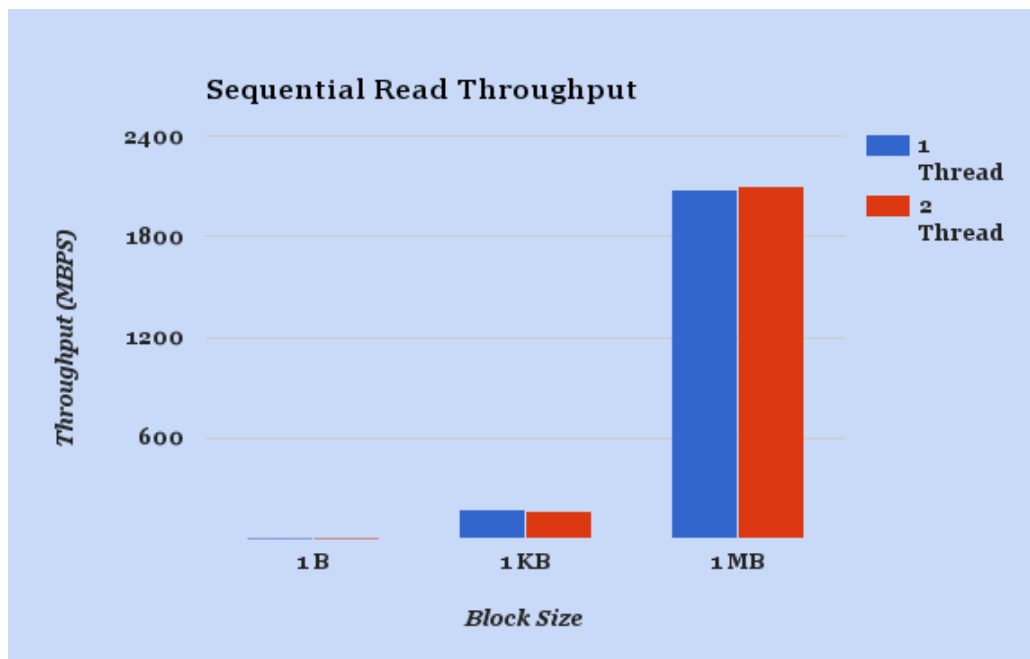
2. DISK BENCHMARKING

Amazon provides with Disk Space of around 8GB and Cache of size 25600Kbytes. The disk benchmarking is evaluated on the basis of these specifications.

The graphs for Throughput and Latency for, Sequential Read & Write as well as Random Read & Write are shown below:

Throughput:

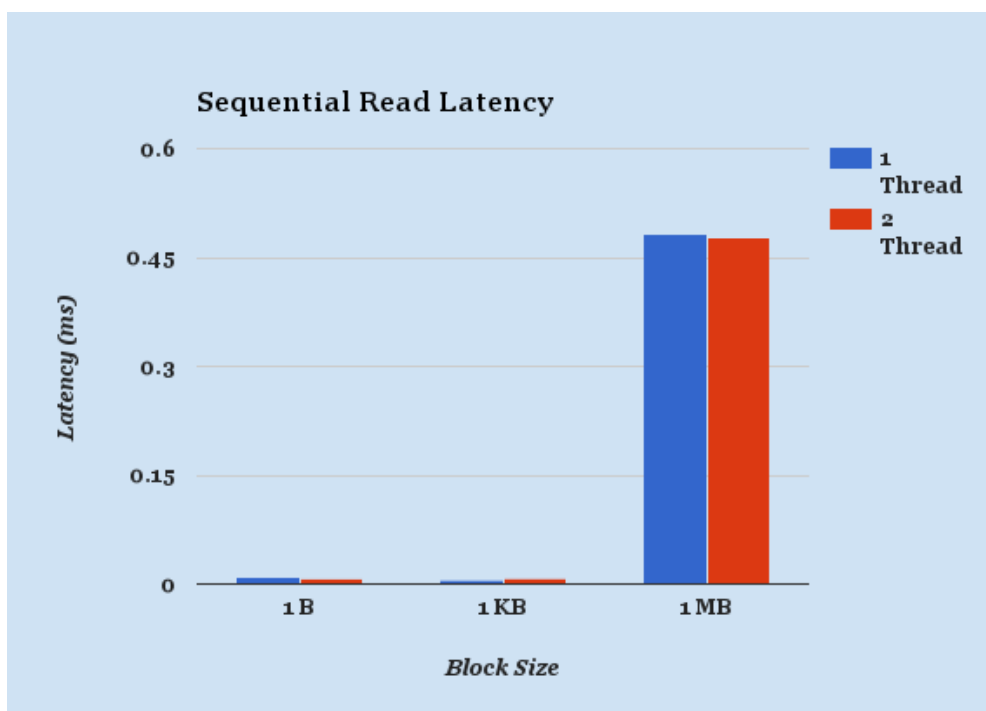
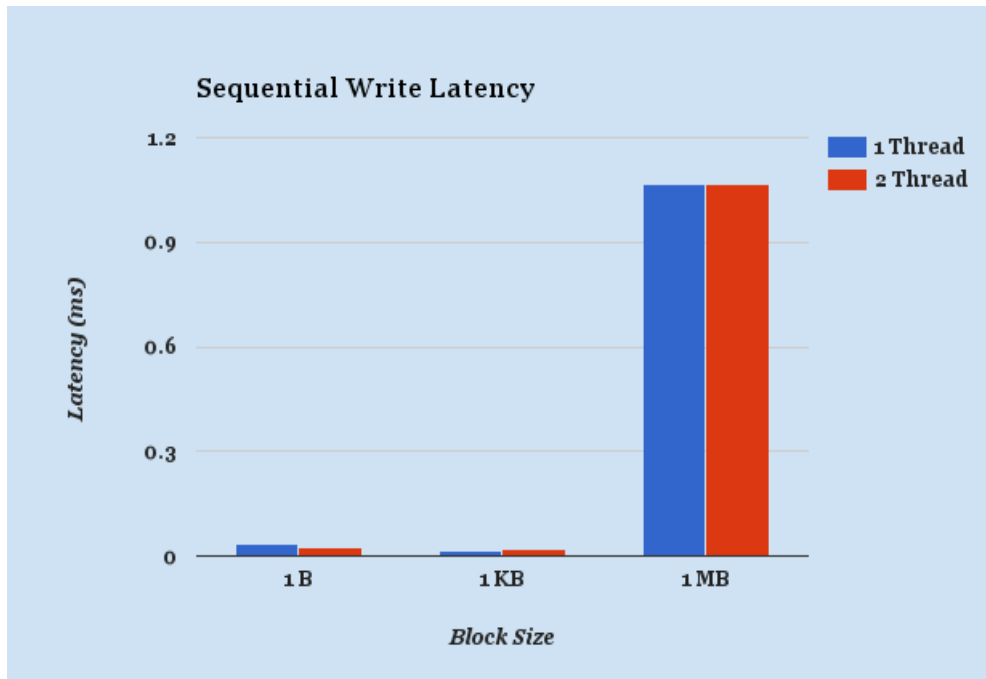


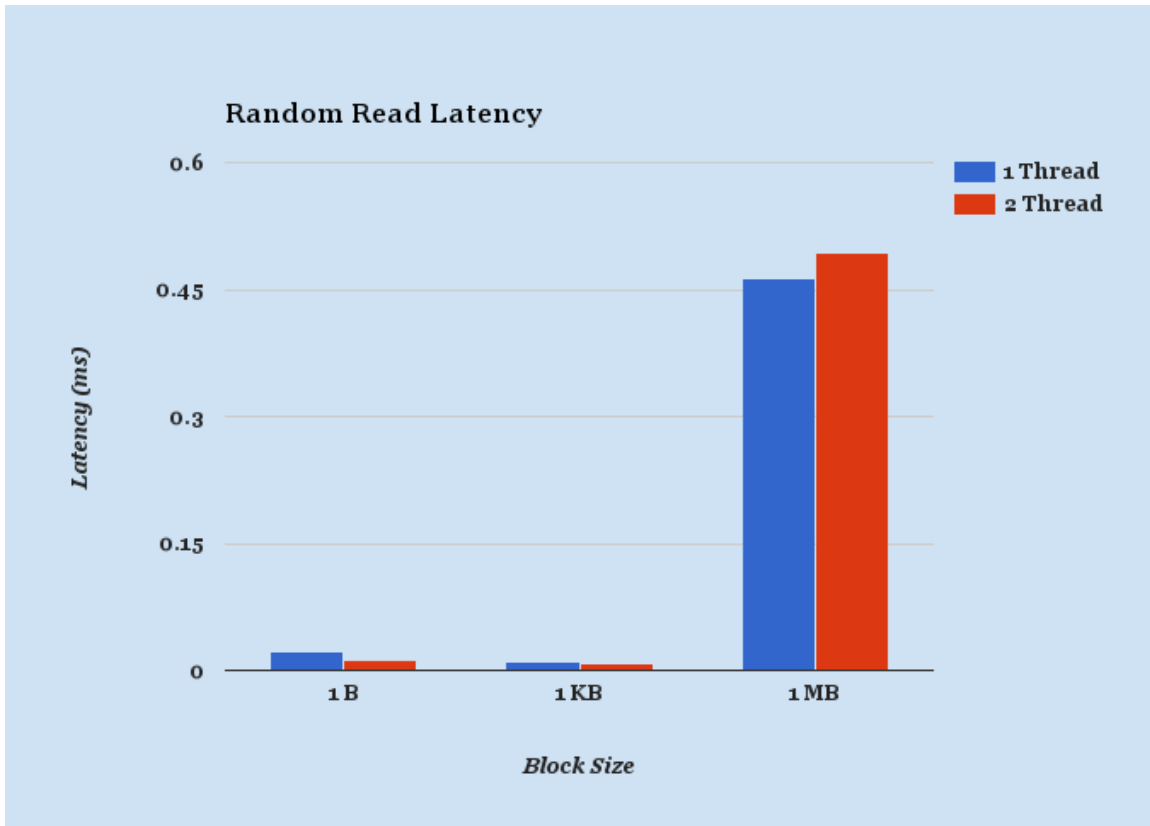
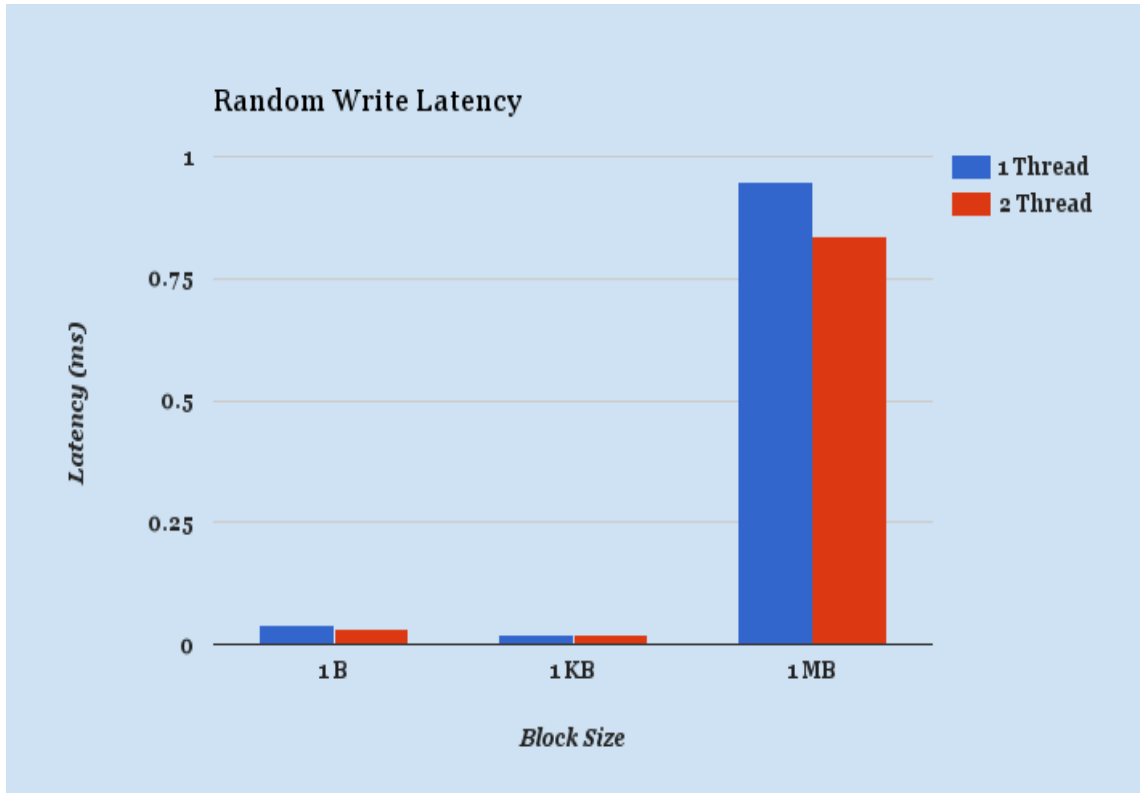


As we increase the number of threads, we can see the Throughput increases; the speeds for Sequential are a bit higher for some cases than Random disk access.

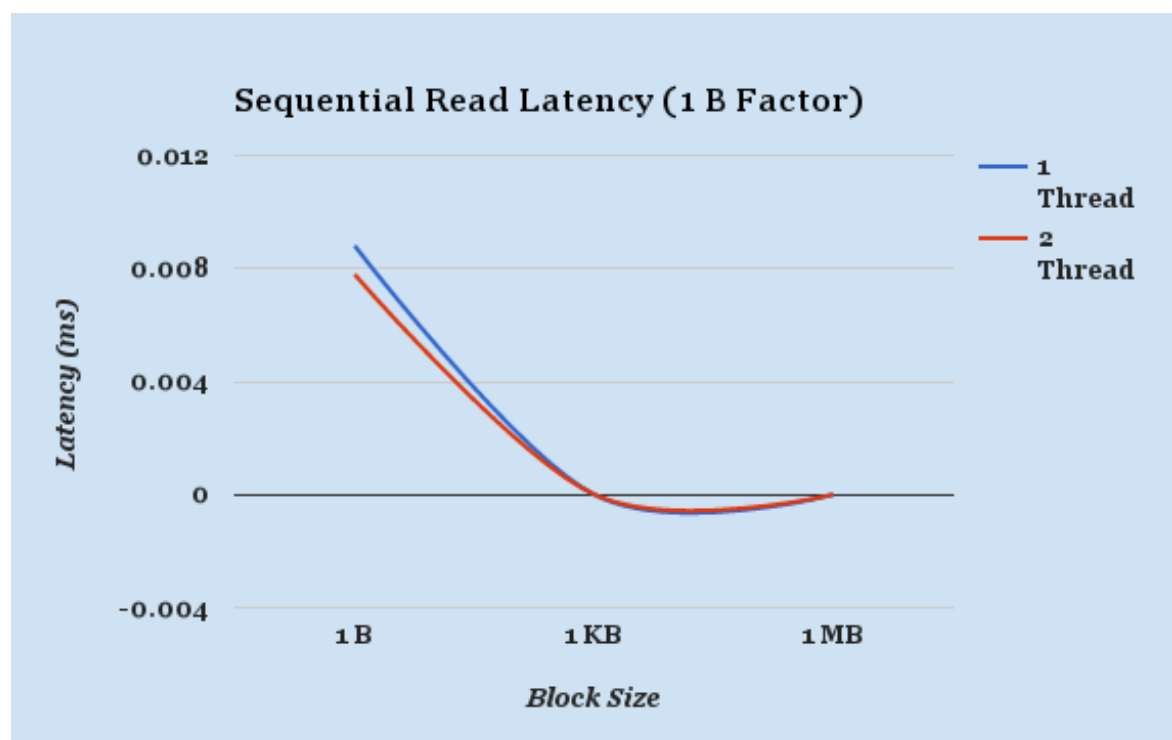
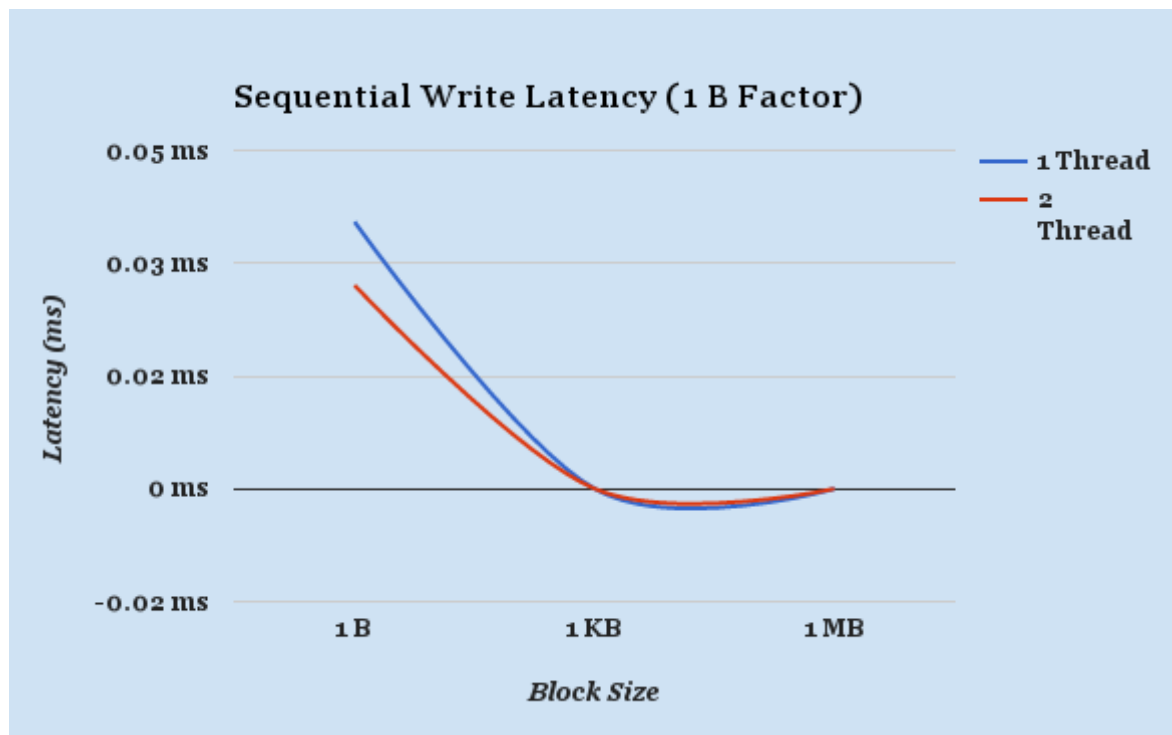
Latency:



As we see the graphs below the latency decreases as we increase the number of threads, as the data to access from disk also increases with the increase in number of threads. The Bar graphs shows and increase in latency, this is because the latencies are calculated for 1B, 1 KB and 1 MB as a whole. The other line graph shows the latencies with common 1Byte factor, for all block sizes.





These latencies graph are considered with the 1 Byte factor, i.e. all the values are for 1Byte of data transfer, from different Block sizes.





The Average and Standard Deviation for 3 Experiments:

Average (Latency in ms):

Block Size	1 B	1KB	1MB
Seq. Write	0.035500667	0.015134667	1.067327
Seq. Read	0.008803	0.005443333	0.481972667
Ran. Write	0.03964	0.018343	0.94679
Ran. Read	0.02223	0.0093596	0.462244

Average (Throughput in MB/s):

Block Size	1 B	1KB	1MB
Seq. Write	0.0269241	64.54552182	934.7162732
Seq. Read	0.108385	179.455102	2077.08631
Ran. Write	0.0240529	53.25690	1066.74833
Ran. Read	0.0463590	118.6682	2161.63714

Standard Deviation (Latency):

Block Size	Seq. Write	Seq. Read	Ran. Write	Ran. Read
1 B	0.002081	0.0002344	0.0002344	0.0002173
1 KB	0.000332	0.0001107	0.0001107	0.0004130
1 MB	0.015442	0.0195914	0.0195914	0.1190715

Standard Deviation (Throughput):

Block Size	Seq. Write	Seq. Read	Ran. Write	Ran. Read
1 B	0.0015896	0.0028439	0.0001315	0.01407758
1 KB	1.4087563	3.6764208	1.1944984	1.74758713
1 MB	12.644688	52.153752	12.004160	52.1278257

IOZone Benchmarking:

The following screen shots shows the performance of IOZone Benchmarking tool on AWS. The result shows the different Speeds in Kbytes. The 1st screen shows the speeds of all the values which gets run by default when executing IOZone. The 2nd screen shows the values for specific file size provided by us.

```
Run began: Sat Feb 13 01:45:55 2016

Auto Mode
Command line used: ./iozone -a
Output is in kBytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 kBytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.

kB reclen write rewrite read reread random read random write read rewrite read fwrite frewrite fread freread
64 4 1279447 3541098 10821524 12902017 7100397 3701156 5860307 4564706 10821524 3541098 4018152 6421025 7940539
64 8 1679761 4018152 15972885 15972885 12902017 4564706 7100397 5283570 22737791 4897948 4274062 9318832 9006179
64 16 2006158 4564706 15972885 15972885 15972885 4897948 7100397 5389653 10821524 3057153 4564706 7940539 15972885
64 32 2006158 4274062 15972885 12902017 12902017 4274062 6271021 4897948 15972885 4643754 6421025 9318832 9006179
64 64 1933893 4564706 20062191 20062191 15972885 4988978 9318832 5389653 9006179 4564706 8182586 15972885
128 4 1642783 4012317 11720614 11720614 9795896 4012317 7082197 5122535 7476717 3759450 4012317 9129573 10779307
128 8 2175277 2066432 16365173 15881078 12542043 4934216 9129573 2409592 5847904 4596273 4557257 10567140 10779307
128 16 2123084 3759450 6114306 14200704 14200704 5122535 8548124 6114306 12842051 3759450 4596273 4596273 6114306
128 32 1111081 3867787 5122535 15881078 15881078 5545860 9977956 6406138 10567140 5122535 5847904 11470204 14200704
128 64 1142750 2326073 18012359 18012359 15881078 5847904 9977956 5847904 5325709 3199360 3895854 12842051 14586625
128 128 1102844 2367096 7082197 7082197 4717434 5603747 9795896 5545860 4934216 5545860 5784891 4934216 14200704
256 4 1842366 4132864 11569783 12812277 10245071 4332998 8888172 4477549 8534922 3935921 4197489 9688433 10245071
256 8 2065659 4734192 14164395 15835569 12812277 4496299 9518507 6719046 12090911 4840911 4819184 12228612 12090911
256 16 2065128 4019184 15164624 14953435 13454450 5687020 10651588 6554072 15164624 3994490 4819184 12090911 12812277
256 32 2639446 5320671 14953435 14164395 14164395 5910112 11091721 6554072 15164624 5217259 5217259 11695808 13454450
256 64 2408631 5540300 14953435 16072608 15164624 5971670 11207494 6719046 13454450 5117791 6213578 13454450 15164624
256 128 2557711 5455847 12812277 17096249 14164395 6107548 11091721 6595234 12812277 4197489 7735574 5117791 14164395
256 256 2632974 4132864 13454450 13454450 12812277 5455847 10245071 5687020 11207494 5217259 5569035 11207494 12661199
512 4 1889050 4195896 11620224 11943357 9510316 4195896 9304292 5566231 8528336 4131319 4091959 9659630 11374040
512 8 2370799 4522868 7540170 15471149 13438092 5278892 11620224 7410080 11620224 4838792 4375425 13109944 14240069
512 16 2141475 5331314 14628067 15883417 13872122 5760329 11943357 6319739 12797441 5019764 4961773 13109944 14628067
512 32 2695115 5699180 14240069 15037801 14240069 6086873 12797441 6999490 15471149 5115422 5164632 13109944 14628067
512 64 2548017 5807059 15037801 15037801 14628067 6174377 12797441 7022379 14146265 5164632 5566231 13438092 15037801
512 128 2815243 5566231 13783088 14628067 13872122 6104175 11877300 6632013 12499490 5566231 6035551 9638369 14240069
512 256 2532990 5398323 12499490 12499490 12146009 5684095 11080601 6246213 11620224 4228947 7309196 11138071 12797441
512 512 2524058 5331314 13109944 12797441 11374040 5684095 11620224 5495016 7758090 3710197 5951911 11620224 11877300
1024 4 1858644 4178773 11773301 12492426 9464330 3941039 9855234 5917516 9677584 3891053 4162573 10769573 10134284
1024 8 2437821 5120336 15295157 16037624 12348754 5018624 13142265 6963241 12037272 5048117 4228138 12790037 14183905
1024 16 2326879 5157871 13472053 15080341 13818817 5277632 11773301 6963241 12639479 4493556 5018624 13264026 14183905
1024 32 2744728 5564829 13472053 14618393 14044758 5660167 13472053 6569179 12492426 4920874 4854136 13643232 14668318
1024 64 2694787 5958564 12348754 15516181 13863422 6025439 13102174 6609617 15080341 5018624 5120336 14044758 15295157
1024 128 2835325 5720477 13643232 14044758 12790037 5917516 8822754 6355328 13818817 5096034 5330028 9485232 13818817
1024 256 2437821 4805258 12037272 12790037 11903823 5720477 8474583 6393168 12173747 4589593 6093831 11368190 13142265
1024 512 2688041 5652717 12944224 12790037 12492426 5690162 11903823 5885083 11773301 3863055 7160597 11368190 12944224
1024 1024 2553783 4972145 11368190 11520658 13142265 5277632 11773301 6025439 12456195 5251818 5782087 12173747 12492426
2048 4 1885905 4221501 11706006 7207495 9062970 4130162 7556202 5754221 9101380 3743207 4138120 9796850 11260973
2048 8 2062505 5119743 13836752 12959958 12118884 5549749 6873016 6895083 11440955 4764758 3483621 7207495 11770166
2048 16 2060526 5131978 12580349 6990474 12805399 5319484 12327589 6523312 6298503 4900677 4552633 13489908 14539356
2048 32 2488533 5277002 12882215 14944066 13926483 5970183 10240672 6807652 9308497 5277002 5135046 13467949 7189398
```

```
ubuntu@ip-172-31-23-104:~/iozone3_434/src/current$ ./iozone -g# -s 1024
Iozone: Performance Test of File I/O
Version $Revision: 3.434 $
Compiled for 64 bit mode.
Build: linux

Contributors: William Norcott, Don Capps, Isom Crawford, Kirby Collins
Al Slater, Scott Rhine, Mike Wisner, Ken Goss
Steve Landherr, Brad Smith, Mark Kelly, Dr. Alain CYR,
Randy Dunlap, Mark Montague, Dan Million, Gavin Brebner,
Jean-Marc Zucconi, Jeff Blomberg, Benny Halevy, Dave Boone,
Erik Habbinga, Kris Strecker, Walter Wong, Joshua Root,
Fabrice Bacchella, Zhenghua Xue, Qin Li, Darren Sawyer,
Vangel Bojaxhi, Ben England, Vikentsi Lapa,
Alexey Skidanov.

Run began: Sat Feb 13 01:49:37 2016

Using maximum file size of 4 kilobytes.
File size set to 1024 kB
Command line used: ./iozone -g# -s 1024
Output is in kBytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 kBytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.

kB reclen write rewrite read reread random read random write read rewrite read fwrite frewrite fread freread
1024 4 1577984 4110780 11398360 11903823 8542002 4392454 8457895 5950309 8262639 4044965 3970183 10429596 9743447

iozone test complete.
```

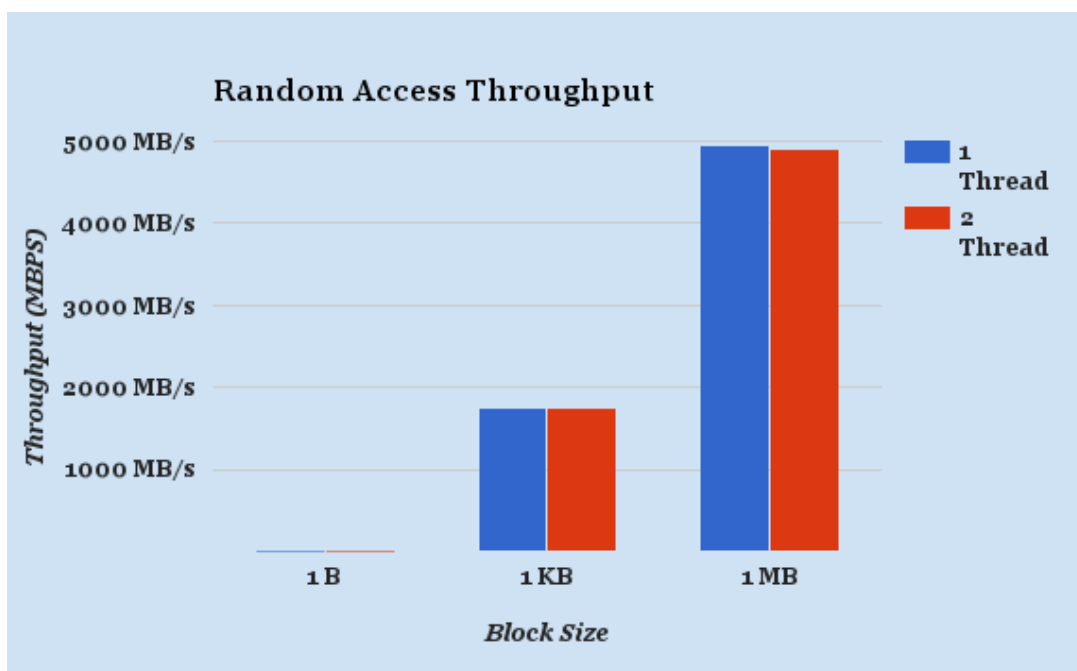
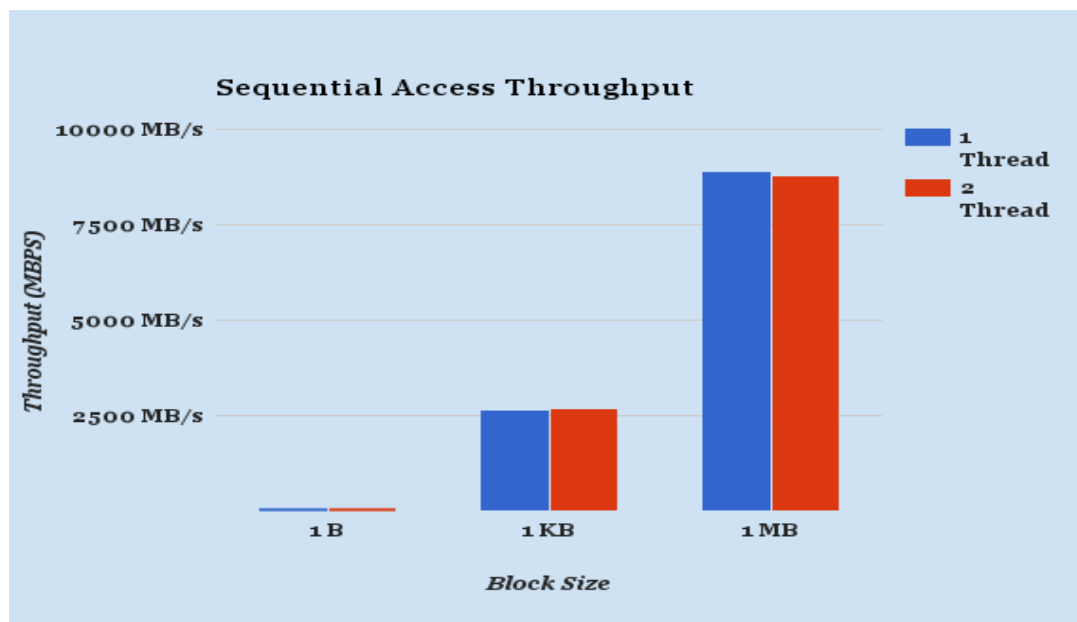
As seen above, the performance for 1MB (1024KB) of file for iozone benchmarking tool is:
Sequential Write: 1.5048 Gbps, Sequential Read: 10.87 Gbps, Random Write: 4.188 Gbps, Random Read: 8.146 Gbps
The efficiency of the tool when compared with theoretical performance for sequential write is about: **51%**.



3. MEMORY BENCHMARKING

Amazon for memory doesn't specify its Base DRAM Clock frequency, which is essential for calculating the Theoretical performance. The Data transfer per clock is 2, Memory Bus width being 64 and Number of Interfaces are 2.

Theoretical Performance (considering my system): Base DRAM Clock freq. * No. of Data transfer per clock * Number of Interfaces = $1600 * 2 * 64 * 2 = 51.2 \text{ GB/s}$

The following graph shows the Throughput and Latency for Sequential and Random access of memory:





As, we can see the latency starts to decrease as we increase the number of threads, at some point of time, the memory speed becomes too fast that the latency tends to become towards power of E -15 (minus 15).

The Average and Standard Deviation for 3 Experiments:

Average (Latency in ms):

Block Size	1 B	1KB	1MB
Seq. Access	1.050E-05	1.3E-09	2.215E-10
Ran. Access	0.000113766	5.41167E-07	1.92233E-07

Average (Throughput in MB/s):

Block Size	1 B	1KB	1MB
Seq. Access	90.796756	2654.4614	8906.71
Ran. Access	8.38496	1762.3905	4931.50911

Standard Deviation (Latency):

Block Size	Seq. Access	Ran. Access
1 B	1062E-07	2.233E-06
1 KB	1.6462E-09	6.997E-09
1 MB	1.0332E-11	2.0793E-09

Standard Deviation (Throughput):

Block Size	Seq. Access	Ran. Access
1 B	0.91697326	0.16612
1 KB	152.3571	22.7043
1 MB	121.122	53.0298

STREAM Benchmarking:

The following screen shots of STREAM benchmarking tool shows the results for Memory Speeds achieved when evaluated on AWS.



The result is divided into 4 parts, having a total of 48.75 GB/s, so the performance for it when compared with theoretical performance shows efficiency of: 95.21 %

```
ubuntu@ip-172-31-23-104:~/iozone3_434/src/current$ gcc -o stream stream.c
ubuntu@ip-172-31-23-104:~/iozone3_434/src/current$ ./stream
-----
STREAM version $Revision: 5.9 $
-----
This system uses 8 bytes per DOUBLE PRECISION word.
-----
Array size = 2000000, Offset = 0
Total memory required = 45.8 MB.
Each test is run 10 times, but only
the *best* time for each is used.
-----
Printing one line per active thread....
-----
Your clock granularity/precision appears to be 1 microseconds.
Each test below will take on the order of 2099 microseconds.
 (= 2099 clock ticks)
Increase the size of the arrays if this shows that
you are not getting at least 20 clock ticks per test.
-----
WARNING -- The above is only a rough guideline.
For best results, please be sure you know the
precision of your system timer.
-----
Function      Rate (MB/s)    Avg time     Min time     Max time
Copy:         12052.5977    0.0028      0.0027      0.0040
Scale:        12158.5042    0.0027      0.0026      0.0028
Add:          13043.5110    0.0037      0.0037      0.0038
Triad:        12671.6133    0.0042      0.0038      0.0075
-----
Solution Validates
-----
```

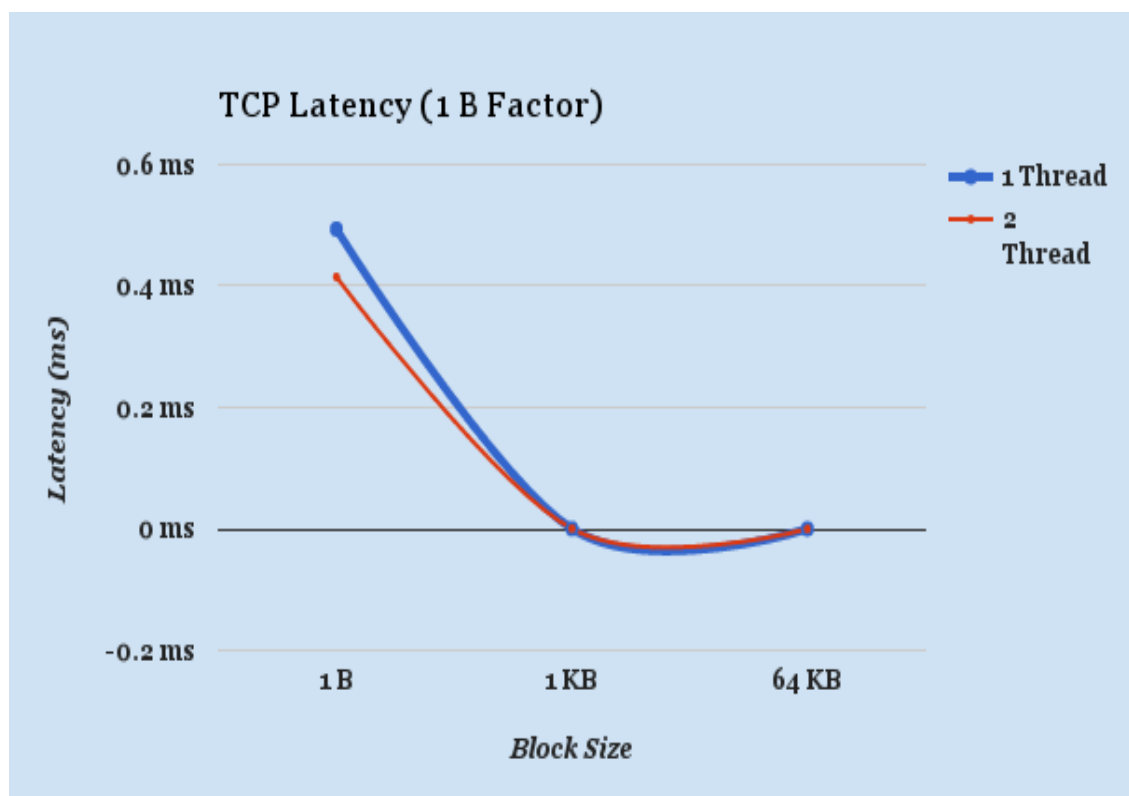
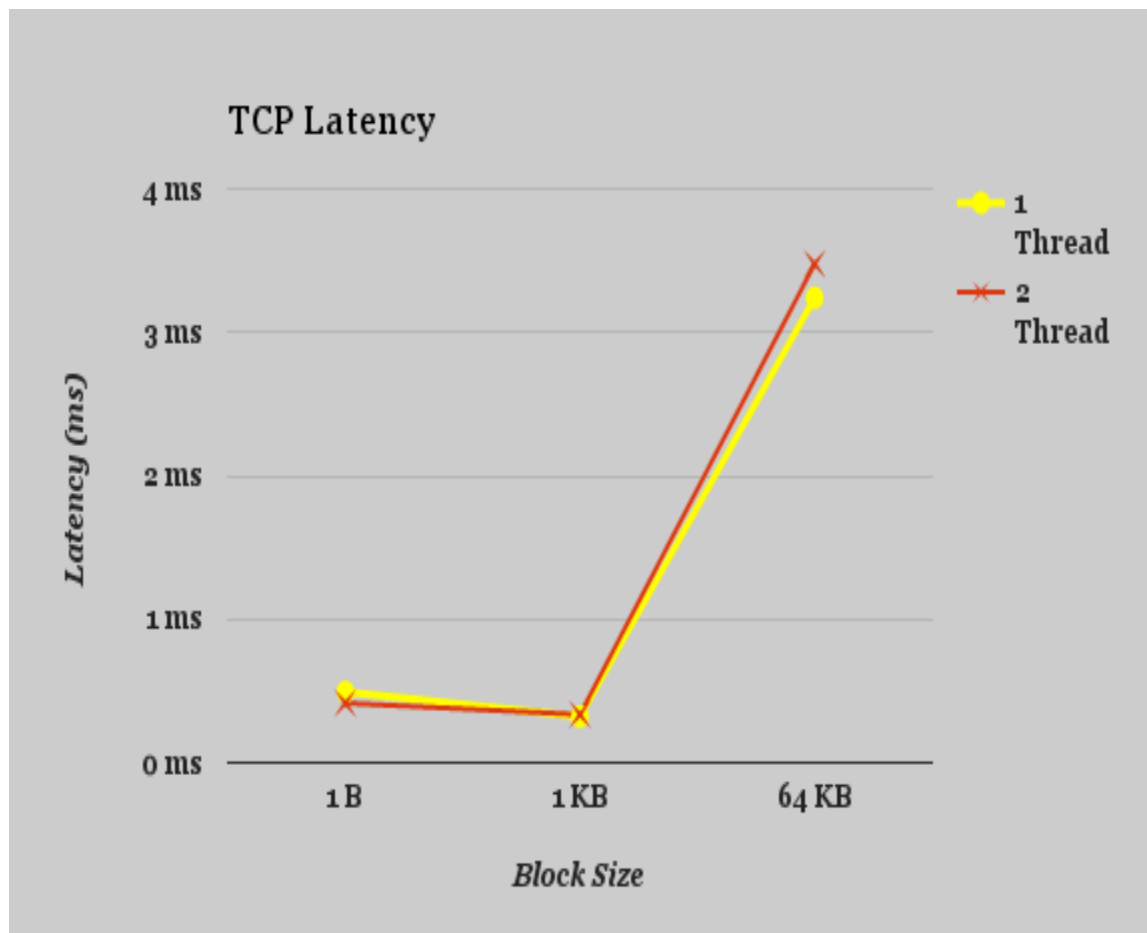
4. NETWORK BENCHMARK



The network benchmark evaluation is performed for TCP as well as UDP. Different Packet sizes are transferred ranging from 1B, 1KB, 64KB and the RTT is calculated accordingly.

The throughput and latency for each of them is shown below in the graphs:



As seen, the network speed i.e. throughput increases with the increase in Packet size to sent for the TCP and UDP network protocol. UDP being connection less protocol is faster when compare with TCP network. There might be a lot of data which is lost when transmitted for UDP.





The above graphs are for Latency, the 1B factor graphs are for comparing all the Packet sizes with the same 1Byte factor for latency. The other two graphs are for latency as a whole, i.e. the entire packet size latency is considered. This is the reason

why the other two graph shows that latency is in increasing order. In actual, the 1 Byte factor graphs tells the real numbers when comparing latency. We observe that latency decreases with the increase in Packet Size to sent across network for both the cases of TCP and UDP.

The Average and Standard Deviation for 3 Experiments:

Average (Latency in ms):

Block Size	1 B	1KB	64KB
TCP	0.493708	0.3216916	3.240265
UDP	0.5723611	0.391104	2.588017

Average (Throughput in MB/s):

Block Size	1 B	1KB	64KB
TCP	0.0020187	3.0817792	19.3021101
UDP	0.001669	2.5429939	24.2412297

Standard Deviation (Latency):

Block Size	TCP	UDP
1 B	0.134623	0.03145243
1 KB	0.050008	0.06583269
64KB	0.135084	0.16675977

Standard Deviation (Throughput):

Block Size	TCP	UDP
1 B	0.00047992	8.89916E-05
1 KB	0.44158978	0.41199362
64KB	0.8165913	1.58918945

IPERF BENCHMARKING: TCP & UDP:

```
ubuntu@ip-172-31-25-167:~$ iperf -c 52.26.161.78 -n 1
-----
Client connecting to 52.26.161.78, TCP port 5001
TCP window size: 325 KByte (default)
-----
[ 3] local 172.31.25.167 port 46111 connected with 52.26.161.78 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec  128 KBytes  18.7 Gbits/sec
ubuntu@ip-172-31-25-167:~$ iperf -c 52.26.161.78 -n 64
-----
Client connecting to 52.26.161.78, TCP port 5001
TCP window size: 325 KByte (default)
-----
[ 3] local 172.31.25.167 port 46112 connected with 52.26.161.78 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec  128 KBytes  18.7 Gbits/sec
ubuntu@ip-172-31-25-167:~$ iperf -c 52.26.161.78 -n 1024
-----
Client connecting to 52.26.161.78, TCP port 5001
TCP window size: 325 KByte (default)
-----
[ 3] local 172.31.25.167 port 46113 connected with 52.26.161.78 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec  128 KBytes  15.7 Gbits/sec
ubuntu@ip-172-31-25-167:~$ iperf -c 52.26.161.78 -n 65536
-----
Client connecting to 52.26.161.78, TCP port 5001
TCP window size: 325 KByte (default)
-----
[ 3] local 172.31.25.167 port 46114 connected with 52.26.161.78 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec  128 KBytes  16.4 Gbits/sec
```

```
error: No address associated with hostname
ubuntu@ip-172-31-25-167:~$ iperf -c 52.26.161.78 -u -n 1
-----
Client connecting to 52.26.161.78, UDP port 5001
Sending 1470 byte datagrams
UDP buffer size: 208 KByte (default)
-----
[ 3] local 172.31.25.167 port 50114 connected with 52.26.161.78 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec  1.44 KBytes  1.04 Mbits/sec
[ 3] Sent 1 datagrams
[ 3] Server Report:
[ 3] 0.0- 0.0 sec  1.44 KBytes  1.04 Mbits/sec  0.000 ms  0/ 1 (0%)
ubuntu@ip-172-31-25-167:~$ iperf -c 52.26.161.78 -u -n 1024
-----
Client connecting to 52.26.161.78, UDP port 5001
Sending 1470 byte datagrams
UDP buffer size: 208 KByte (default)
-----
[ 3] local 172.31.25.167 port 56925 connected with 52.26.161.78 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec  1.44 KBytes  1.04 Mbits/sec
[ 3] Sent 1 datagrams
[ 3] Server Report:
[ 3] 0.0- 0.0 sec  1.44 KBytes  1.05 Mbits/sec  0.000 ms  0/ 1 (0%)
ubuntu@ip-172-31-25-167:~$ iperf -c 52.26.161.78 -u -n 65536
-----
Client connecting to 52.26.161.78, UDP port 5001
Sending 1470 byte datagrams
UDP buffer size: 208 KByte (default)
-----
[ 3] local 172.31.25.167 port 55050 connected with 52.26.161.78 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.5 sec  64.6 KBytes  1.05 Mbits/sec
[ 3] Sent 45 datagrams
[ 3] Server Report:
[ 3] 0.0- 0.5 sec  64.6 KBytes  1.05 Mbits/sec  0.041 ms  0/ 45 (0%)
```